
THE MILLION-LABEL NER: BREAKING SCALE BARRIERS WITH GLINER BI-ENCODER

Ihor Stepanov¹, Mykhailo Shtopko¹, Dmytro Vodianytskyi¹, Oleksandr Lukashov¹

¹Knowledgator Engineering, Kyiv, Ukraine

Correspondence: ingvarstep@knowledgator.com, mykhailoshtopko@knowledgator.com

ABSTRACT

This paper introduces GLiNER-bi-Encoder, a novel architecture for Named Entity Recognition (NER) that harmonizes zero-shot flexibility with industrial-scale efficiency. While the original GLiNER framework offers strong generalization, its joint-encoding approach suffers from quadratic complexity as the number of entity labels increases. Our proposed bi-encoder design decouples the process into a dedicated label encoder and a context encoder, effectively removing the context-window bottleneck. This architecture enables the simultaneous recognition of thousands—and potentially millions—of entity types with minimal overhead. Experimental results demonstrate state-of-the-art zero-shot performance, achieving 61.5% Micro-F1 on the CrossNER benchmark. Crucially, by leveraging pre-computed label embeddings, GLiNER-bi-Encoder achieves up to a 130× throughput improvement at 1024 labels compared to its uni-encoder predecessors. Furthermore, we introduce GLiNTER, a modular framework that leverages this architecture for high-performance entity linking across massive knowledge bases such as Wikidata.

Keywords GLiNER · Information Extraction · NLP · NER · Zero-shot classification · BERT · ModernBERT

1 Introduction

Information extraction (IE) remains a fundamental challenge in natural language processing, with critical applications spanning scientific research [Hong et al., 2021], healthcare [Yazdani et al., 2024], finance [Skalický et al., 2022], and public administration [Siciliani et al., 2023]. The exponential growth of unstructured text data across these domains demands IE systems that can efficiently process large volumes while maintaining high accuracy—particularly in high-stakes applications where extraction errors can have serious consequences. Moreover, the dynamic nature of real-world applications requires models that can rapidly adapt to new entity types and domains without extensive retraining.

Named Entity Recognition (NER), as a cornerstone IE task, has evolved significantly from early rule-based systems to modern neural approaches. While rule-based methods [Petasis et al., 2001] required extensive manual effort and lacked generalization capabilities, statistical approaches like Hidden Markov Models [Seymore et al., 1999] and Conditional Random Fields [McDonald and Pereira, 2005] introduced data-driven learning but struggled with complex contextual dependencies. The transformer revolution, initiated by BERT [Devlin et al., 2018], transformed NER by enabling bidirectional context understanding, leading to substantial performance improvements across benchmarks.

Recent advances have explored two primary paradigms for zero-shot and few-shot NER. Generative approaches, exemplified by large language models like GPT-4 [OpenAI et al., 2024], frame NER as a text generation task, offering flexibility in handling diverse entity types through natural language prompts. However, these methods suffer from computational inefficiency—generating tokens for entities already present in the text—and often struggle with structured output consistency, leading to hallucinations and format violations that compromise reliability in production systems.

In contrast, discriminative approaches have shown promise in maintaining efficiency while achieving competitive performance. GLiNER [Zaratiana et al., 2023] introduced a unified token classification architecture that encodes both text and entity type descriptions via a single transformer, enabling zero-shot entity recognition through attention-based label-text interactions. While effective, this approach faces scalability challenges when handling large label sets, as all

labels must be processed jointly with each input text, resulting in quadratic complexity with respect to the number of entity types.

In this work, we present a bi-encoder architecture that addresses these limitations by decoupling text and entity type encodings. Our approach separates the encoding process into two specialized components: a text encoder that processes input sequences to generate span representations, and a label encoder that independently embeds entity type descriptions. This separation offers several key advantages: (1) entity type embeddings can be pre-computed and cached, significantly reducing inference time for large label sets; (2) leveraging pre-trained sentence transformers as label encoders provides superior semantic understanding for unseen entity types; and crucially, (3) the decoupled design enables the model to efficiently handle thousands or even millions of entity types simultaneously, as label embeddings are computed independently and can be stored in efficient vector databases for similarity search.

This scalability is particularly important for real-world applications such as biomedical entity linking, where ontologies like UMLS [Bodenreider, 2004] contain over 4 million concepts, or enterprise knowledge extraction systems that must recognize entities across numerous taxonomies. Unlike joint encoding approaches that face memory and computational constraints when scaling beyond hundreds of labels, our bi-encoder architecture maintains constant memory usage for text encoding regardless of the number of entity types, with label retrieval complexity reduced to an efficient nearest-neighbor search.

Experimental results demonstrate that our bi-encoder approach achieves state-of-the-art performance on zero-shot NER benchmarks while offering 2-3x faster inference than uni-encoder GLiNER with dozens of entity types, with the performance gap widening to 10-100x when scaling to thousands of entity types. Moreover, the model exhibits strong generalization to specialized domains without fine-tuning, making it particularly suitable for real-world deployment scenarios where entity types evolve dynamically, and scale requirements are demanding.

Beyond standalone named entity recognition, the bi-encoder architecture’s ability to efficiently handle large entity type sets makes it particularly well-suited for entity linking tasks. Entity linking extends NER by disambiguating mentions to specific knowledge base entities, requiring candidate retrieval and entity disambiguation—operations that benefit directly from precomputed embeddings and efficient similarity search. To facilitate practical deployment of GLiNER-bi-Encoder in entity linking scenarios, we have developed GLiNKER, a modular DAG-based framework that integrates bi-encoder models into multi-stage linking pipelines combining mention extraction, candidate generation from knowledge bases, and linking entities.

2 Methods

Zero-shot NER requires computing compatibility between text spans and arbitrary entity type definitions. The original GLiNER architecture addresses this by jointly encoding entity types with the input text and processing them with a single transformer. While this enables rich cross-attention between labels and context, it couples computational cost to label count—processing n tokens against m entity types requires self-attention over sequences of length $\mathcal{O}(n + m)$, with complexity $\mathcal{O}((n + m)^2)$.

The bi-encoder architecture restructures this computation by observing that entity type semantics are largely independent of input text. A description like “person” or “chemical compound” can be encoded once and reused across all inputs. This enables architectural separation: independent encoders process text and labels in parallel, combining embeddings only at the final scoring stage. Text encoding becomes $\mathcal{O}(n^2)$ independent of label count, while label encoding can be performed offline and cached. With pre-computed embeddings, inference complexity is effectively $\mathcal{O}(n^2)$ regardless of whether the system recognizes 10 or 10,000 entity types.

This separation also enables heterogeneous encoder selection—pairing text encoders optimized for contextual understanding with label encoders optimized for semantic similarity. To address the potential loss of cross-modal interaction, an optional cross-attention fusion layer can exchange information between representations after initial encoding, offering a tunable trade-off between accuracy and efficiency.

2.1 Architecture Overview

The bi-encoder architecture employs two specialized, independent transformers. This separation is the fundamental driver of the model’s scalability, as it allows for heterogeneous backbone selection (e.g., pairing a fast text encoder with a semantically rich sentence-transformer for labels).




Figure 1: GLiNER architecture variants. Left: Uni-encoder processes entity types and text jointly. Right: Bi-encoder uses separate encoders with optional cross-attention fusion.

2.2 Uni-Encoder Architecture

Before jumping into the details of the bi-encoder architecture, let's discuss the essence of the original GLiNER architecture. The uni-encoder processes entity types and input text within a single transformer. Entity types are prepended to the input using special tokens:

$$\mathbf{X} = [[\text{ENT}] t_1 [\text{ENT}] t_2 \dots [\text{SEP}] w_1 w_2 \dots w_L] \quad (1)$$

where t_i denotes entity type tokens and w_j denotes input words. The transformer produces contextualized embeddings:

$$\mathbf{H} = \text{Transformer}(\mathbf{X}) \in \mathbb{R}^{N \times D} \quad (2)$$

Entity type embeddings $\mathbf{E} \in \mathbb{R}^{C \times D}$ are extracted at $[\text{ENT}]$ token positions, and word embeddings $\mathbf{W} \in \mathbb{R}^{L \times D}$ are extracted using a word-to-subtoken mapping that selects the first subtoken of each word.

2.3 Bi-Encoder Architecture

The bi-encoder employs separate transformers for entity types and input text:

$$\mathbf{E} = \text{LabelEncoder}(\mathbf{T}) \in \mathbb{R}^{C \times D} \quad (3)$$

$$\mathbf{H} = \text{TextEncoder}(\mathbf{X}) \in \mathbb{R}^{N \times D} \quad (4)$$

where \mathbf{T} contains tokenized entity type descriptions. Word embeddings \mathbf{W} are extracted from \mathbf{H} as in the uni-encoder case. The label embeddings are expanded across the batch: $\mathbf{E}_{\text{batch}} = \text{expand}(\mathbf{E}, B) \in \mathbb{R}^{B \times C \times D}$.

This design enables caching of entity type representations, significantly improving inference efficiency when the same entity types are queried repeatedly.

2.3.1 Cross-Attention Fusion

Both architectures optionally incorporate bidirectional cross-attention to exchange information between text and entity representations:

$$\mathbf{W}', \mathbf{E}' = \text{CrossFuser}(\mathbf{W}, \mathbf{E}, \mathbf{M}_{\text{text}}, \mathbf{M}_{\text{labels}}) \quad (5)$$

The CrossFuser applies multi-head attention in both directions according to a configurable schema, allowing entity-aware text representations and text-aware entity representations. This fusion occurs after initial encoding and before span/token scoring.

2.4 Span-Level Prediction

Span-level models enumerate all candidate spans up to a maximum width K and score each against entity types.

2.4.1 Span Representation

Word embeddings are projected through separate MLPs for span boundaries:

$$\mathbf{H}_{\text{start}} = \text{MLP}_{\text{start}}(\mathbf{W}) \in \mathbb{R}^{B \times L \times D} \quad (6)$$

$$\mathbf{H}_{\text{end}} = \text{MLP}_{\text{end}}(\mathbf{W}) \in \mathbb{R}^{B \times L \times D} \quad (7)$$

For each span $(i, i + k)$ where $k \in [0, K - 1]$:

$$\mathbf{s}_{i,k} = \text{MLP}_{\text{out}} (\text{ReLU}([\mathbf{h}_{\text{start}}^i \| \mathbf{h}_{\text{end}}^{i+k}])) \quad (8)$$

yielding span representations $\mathbf{S} \in \mathbb{R}^{B \times L \times K \times D}$.

2.4.2 Scoring

Entity type embeddings are projected and scored against spans via:

$$\mathbf{O} = \mathbf{S} \cdot \text{MLP}_{\text{prompt}}(\mathbf{E})^\top \in \mathbb{R}^{B \times L \times K \times C} \quad (9)$$

where $O_{b,i,k,c}$ represents the score for span $(i, i + k)$ belonging to entity class c in batch example b .

2.5 Token-Level Prediction

Token-level models classify each word position independently, predicting BIO-style tags for each entity type.

2.5.1 Token Scoring

A scoring layer computes compatibility between word embeddings and entity types:

$$\mathbf{O} = \text{Scorer}(\mathbf{W}, \mathbf{E}) \in \mathbb{R}^{B \times L \times C \times 3} \quad (10)$$

The final dimension encodes [start, end, inside] probabilities for BIO tagging.

2.5.2 Span Extraction from Tokens

Entity spans are extracted by identifying consistent BIO sequences:

1. Identify positions where $P(\text{start}) > \tau$ and $P(\text{end}) > \tau$
2. For each (start, end) pair of the same class where $\text{start} \leq \text{end}$:
3. Validate that all intermediate positions satisfy $P(\text{inside}) > \tau$
4. Accept spans where inside predictions span the full interval

2.5.3 Optional Span Representation

Token-level models can optionally compute explicit span representations for extracted spans:

$$\mathbf{s}_{(i,j)} = \text{SpanRep}(\mathbf{W}, i, j) \quad (11)$$

These representations enable a secondary span-level classification loss:

$$\mathcal{L}_{\text{total}} = \lambda_{\text{token}} \mathcal{L}_{\text{token}} + \lambda_{\text{span}} \mathcal{L}_{\text{span}} \quad (12)$$

2.6 Training

2.6.1 Loss Function

We use focal loss Lin et al. [2018] to handle class imbalance:

$$\mathcal{L}_{\text{focal}} = -\alpha(1 - p_t)^\gamma \log(p_t) \quad (13)$$

where $p_t = \sigma(o)$ for positive examples and $p_t = 1 - \sigma(o)$ for negatives.

2.6.2 Negative Sampling

Given the quadratic number of candidate spans, we apply stochastic negative masking:

$$m_{i,k,c} = \begin{cases} 1 & \text{if } y_{i,k,c} = 1 \\ \text{Bernoulli}(\rho) & \text{if } y_{i,k,c} = 0 \end{cases} \quad (14)$$

where ρ controls the proportion of negative examples retained. Masking strategies include global (uniform), label-wise, and span-wise sampling.

2.7 Inference

2.7.1 Greedy Span Decoding

Given logits \mathbf{O} , we decode entities through threshold filtering and greedy selection:

1. **Filter:** Extract candidates $\mathcal{C} = \{(s, e, c, p) : \sigma(O_{s,k,c}) > \tau\}$ where $e = s + k$
2. **Sort:** Order candidates by probability descending
3. **Select:** Greedily accept spans that don't conflict with already-selected spans
4. **Output:** Return spans sorted by position

2.7.2 Overlap Strategies

The conflict detection function supports multiple NER paradigms:

Flat NER: Two spans conflict if they overlap:

$$\text{conflict}(s_1, e_1, s_2, e_2) = \max(s_1, s_2) \leq \min(e_1, e_2) \quad (15)$$

Multi-label NER: Spans of different types may overlap:

$$\text{conflict}((s_1, e_1, c_1), (s_2, e_2, c_2)) = (c_1 = c_2) \wedge \text{overlap}(s_1, e_1, s_2, e_2) \quad (16)$$

Nested NER: Proper containment is permitted:

$$\text{conflict} = \text{overlap} \wedge \neg\text{contained} \quad (17)$$

where $\text{contained}((s_1, e_1), (s_2, e_2)) = (s_1 \leq s_2 \wedge e_2 \leq e_1) \vee (s_2 \leq s_1 \wedge e_1 \leq e_2)$.

2.8 Data

Pre-training corpus: The pre-training dataset consists of 8M samples from m-a-p/FineFineWeb for Large, Base, and Small variants. The Edge variant was trained on an extended corpus of 10M samples. Each text was annotated with GPT-4o to generate true candidate labels for entity recognition training.

Post-training corpus: The post-training dataset consists of 40k high-quality samples with sequences up to 2048 tokens, designed to refine model performance on longer contexts and improve generalization across diverse entity types.

2.9 Model Training

2.9.1 Model Variants

We trained four model variants with different capacity-efficiency trade-offs based on the ModernBERT encoder architecture:

- **Large:** 400M parameters (ettin-encoder-400m), 28 layers, 16 attention heads, hidden size 1024, paired with BAAI/bge-base-en-v1.5 label encoder (768-dim)
- **Base:** 150M parameters (ettin-encoder-150m), 22 layers, 12 attention heads, hidden size 768, paired with BAAI/bge-small-en-v1.5 label encoder (384-dim)
- **Small:** 68M parameters (ettin-encoder-68m), 19 layers, 8 attention heads, hidden size 512, paired with sentence-transformers/all-MiniLM-L12-v2 label encoder (384-dim)
- **Edge:** 32M parameters (ettin-encoder-32m), 10 layers, 6 attention heads, hidden size 384, paired with sentence-transformers/all-MiniLM-L6-v2 label encoder (384-dim)

All variants share the ModernBERT architecture configuration with position-free embeddings (sans_pos), local attention window of 128 tokens, global attention every 3 layers, RoPE theta of 160,000, maximum position embeddings of 7,999, and vocabulary size of 50,368 tokens.

2.9.2 Training Stages

Pre-Training: Initial training on the pre-training corpus for one complete epoch with the following configuration:

- Dataset size: 8M samples (Large/Base/Small), 10M samples (Edge)
- Maximum sequence length: 1024 tokens
- Focal loss α parameter: 0.7
- Focal loss γ parameter: 2.0
- Training steps: 500k (Large), 250k (Base/Small), 312.5k (Edge)
- Batch size: 16 (Large), 32 (Base/Small/Edge)

Post-Training: Final stage training on the post-training corpus for one complete epoch with increased context length and adjusted loss parameters:

- Maximum sequence length: 2048 tokens
- Focal loss α parameter: 0.8
- Focal loss γ parameter: 2.0
- Loss reduction: sum aggregation

2.9.3 Optimization

We employed the AdamW optimizer with differential learning rates for the encoder and other components:

- Encoder learning rate: 1×10^{-5}
- Other components learning rate: 3×10^{-5}
- Weight decay (encoder): 0.01
- Weight decay (other): 0.01
- Gradient clipping: maximum norm of 10.0
- Learning rate scheduler: cosine annealing
- Warmup ratio: 0.1 (10% of total training steps)

2.9.4 Architectural Configuration

The models were configured with the following entity recognition parameters:

- Span detection mode: MarkerV0
- Subtoken pooling strategy: first token
- Maximum entity types per batch: 100
- Maximum span width: 12 tokens
- Dropout rate: 0.35
- RNN component: enabled

2.9.5 Training Details

Models were trained with the following data processing and regularization settings:

- Word splitter: whitespace tokenization
- Entity type shuffling: enabled during training
- Maximum negative type ratio: 1.0
- Label smoothing: 0.0 (disabled)

All models utilized pre-trained label encoders for encoding entity type, which were fine-tuned jointly with the main encoder during training. The focal loss formulation helps address class imbalance by focusing on hard examples while down-weighting easy negatives, particularly important for NER, where non-entity tokens vastly outnumber entity tokens.

3 Evaluation

The model was compared with other GLiNER-type models. All models were evaluated on cross-domain NER benchmarks in a zero-shot setting. We documented Micro-F1 scores across all datasets. The datasets encompass diverse domains including AI, Literature, Politics, Science, Movies, and other (see Tables 2 and 3). A threshold of 0.4 was applied to filter predicted entities, and all evaluations were conducted at the span level.

Inference speed is measured on a single NVIDIA H100 GPU with batch size 1. We test across label counts $L \in \{1, 2, 4, 8, 16, 32, 64, 128, 256, 512, 1024\}$ and input lengths $T \in \{64, 256, 512\}$ tokens. For each (L, T) configuration, we execute 10 forward passes and report average throughput in examples per second.

3.1 Model Configuration and Performance

We evaluated four bi-encoder model variants ranging from 60M to 530M parameters, utilizing text encoders from the Ettin suite [Weller et al., 2025] paired with specialized label encoders from sentence-transformers and BGE families [Xiao et al., 2023]. Table 1 presents the complete model specifications and performance metrics.

Table 1: GLiNER-bi-V2 Models Overview

Model name	Params	Encoder	Labels Encoder	Avg. CrossNER Benchmark	Inference Speed (H100, examples/s)	Inference Speed (pre-computed labels)
gliner-bi-edge-v2.0	60 M	ettin-encoder-32m	all-MiniLM-L6-v2	54.0%	13.64	24.62
gliner-bi-small-v2.0	108 M	ettin-encoder-68m	all-MiniLM-L12-v2	57.2%	7.99	15.22
gliner-bi-base-v2.0	194 M	ettin-encoder-150m	bge-small-en-v1.5	60.3%	5.91	9.51
gliner-bi-large-v2.0	530 M	ettin-encoder-400m	bge-base-en-v1.5	61.5%	2.68	3.60

The bi-encoder architecture demonstrates consistent performance improvements with model scaling. On the CrossNER benchmark, performance increases from 54.0% (edge) to 61.5% (large), representing a 13.9% relative improvement. Label pre-computation yields 1.34-1.91× speedup across all model sizes, with the most substantial gains observed in smaller models. The base variant (194M parameters) achieves 98% of the large model's performance while operating 2.6× faster, establishing it as the optimal configuration for most deployment scenarios.

3.2 Zero-shot NER Benchmarks

Table 2 presents comprehensive evaluation results across 19 diverse NER datasets spanning biomedical, social media, news, and technical domains. Performance patterns reveal strong domain generalization capabilities, with WikiNeural (76.6-80.0%), bc5cdr (68.5-73.0%), and Broad Tweet Corpus (70.0-72.1%) showing the highest F1 scores across model variants.

Table 2: Comparison of `gliner-bi-v2.0` models on zero-shot NER benchmarks

Dataset	<code>gliner-bi-edge-v2.0</code>	<code>gliner-bi-small-v2.0</code>	<code>gliner-bi-base-v2.0</code>	<code>gliner-bi-large-v2.0</code>
ACE 2004	26.4%	27.5%	28.9%	31.9%
ACE 2005	26.2%	28.1%	30.0%	31.4%
AnatEM	39.1%	43.6%	35.4%	39.5%
Broad Tweet Corpus	70.0%	71.7%	72.1%	70.9%
CoNLL 2003	61.6%	64.2%	65.6%	66.5%
FabNER	22.4%	23.2%	24.3%	22.7%
FindVehicle	35.6%	40.3%	40.6%	39.1%
GENIA_NER	50.1%	53.8%	56.8%	60.1%
HarveyNER	15.0%	10.6%	12.6%	14.7%
MultiNERD	64.6%	66.0%	68.0%	64.0%
Ontonotes	31.4%	31.9%	33.3%	32.5%
PolyglotNER	45.1%	46.3%	46.6%	46.8%
TweetNER7	36.9%	40.9%	40.4%	41.7%
WikiANN_en	52.3%	54.0%	54.9%	56.3%
WikiNeural	78.0%	79.9%	80.0%	76.6%
bc2gm	58.1%	59.9%	62.7%	61.4%
bc4chemd	45.8%	49.1%	53.6%	50.5%
bc5cdr	68.5%	71.5%	73.0%	71.7%
ncbi	65.9%	65.4%	65.2%	65.9%
Average	47.0%	48.8%	49.7%	49.7%
Dataset	<code>gliner-bi-edge-v2.0</code>	<code>gliner-bi-small-v2.0</code>	<code>gliner-bi-base-v2.0</code>	<code>gliner-bi-large-v2.0</code>
CrossNER_AI	53.8%	54.7%	58.3%	57.4%
CrossNER_literature	56.2%	62.6%	65.2%	63.2%
CrossNER_music	68.2%	72.3%	73.4%	74.0%
CrossNER_politics	68.7%	70.0%	70.8%	73.0%
CrossNER_science	63.2%	66.1%	68.0%	67.6%
mit-movie	30.5%	35.2%	46.2%	51.0%
mit-restaurant	37.1%	39.5%	40.3%	44.3%
Average	54.0%	57.2%	60.3%	61.5%

Biomedical benchmarks exhibit strong scaling effects, with GENIA_NER improving from 50.1% to 60.1% (20% relative gain) and bc4chemd from 45.8% to 53.6% between edge and base variants. Performance on social media datasets remains stable across model sizes, suggesting that informal text benefits less from increased model capacity. Challenging domains such as HarveyNER (10.6-15.0%) and FabNER (22.4-24.3%) show limited improvements with scaling, indicating inherent task difficulty rather than model limitations.

3.3 Comparison with Uni-encoder Architecture

Comparative evaluation against uni-encoder GLiNER models (Table 3) reveals comparable performance between architectures on standard benchmarks. The bi-encoder GLiNER-bi-large-v2.0 achieves 61.5% on CrossNER versus 60.9% for `gliner-large-v2.5`, while maintaining substantially different computational characteristics. On the broader benchmark suite, bi-encoder models achieve 49.7% average F1 compared to 46.4-47.0% for uni-encoder variants.

Table 3: Comparison of `gliner-community v2.5` models on zero-shot NER benchmarks

Dataset	<code>gliner_small-v2.5</code>	<code>gliner_medium-v2.5</code>	<code>gliner_large-v2.5</code>
ACE 2004	28.2%	25.9%	26.8%
ACE 2005	28.5%	25.8%	27.2%
AnatEM	39.7%	43.1%	38.1%
Broad Tweet Corpus	65.0%	66.5%	66.7%
CoNLL 2003	65.4%	64.9%	64.2%
FabNER	22.4%	23.5%	25.1%
FindVehicle	39.0%	41.1%	47.0%
GENIA_NER	48.5%	55.1%	46.9%
HarveyNER	17.4%	18.9%	19.1%
MultiNERD	57.0%	59.6%	50.9%
Ontonotes	27.3%	26.1%	22.8%
PolyglotNER	42.0%	43.1%	42.4%
TweetNER7	38.9%	38.6%	38.9%
WikiANN_en	57.3%	58.1%	59.0%
WikiNeural	71.4%	72.6%	73.0%
bc2gm	51.4%	55.6%	52.5%
bc4chemd	38.7%	43.9%	46.8%
bc5cdr	65.9%	65.9%	68.7%
ncbi	64.2%	64.3%	65.3%
Average	45.7%	47.0%	46.4%
Dataset (zero-shot)	<code>gliner_small-v2.5</code>	<code>gliner_medium-v2.5</code>	<code>gliner_large-v2.5</code>
CrossNER_AI	54.2%	54.7%	54.9%
CrossNER_literature	62.5%	66.1%	65.2%
CrossNER_music	68.9%	72.9%	71.6%
CrossNER_politics	64.1%	68.1%	72.4%
CrossNER_science	63.5%	66.3%	63.9%
mit-movie	50.3%	41.4%	53.0%
mit-restaurant	39.8%	38.5%	45.0%
Average	57.6%	58.3%	60.9%

Domain-specific comparisons show bi-encoder advantages on GENIA_NER (60.1% vs 46.9%), MultiNERD (68.0% vs 50.9%), and Ontonotes (33.3% vs 22.8%). Uni-encoder models perform better on CoNLL 2003 (65.4% vs 66.5%) and certain specialized datasets. These patterns suggest that bi-encoder architectures excel when semantic similarity between labels and text is crucial, while uni-encoder models maintain advantages for well-established benchmark formats.

3.4 Inference Speed and Scalability

Table 4 presents inference speed measurements across varying label counts. The bi-encoder architecture demonstrates superior scaling characteristics compared to uni-encoder approaches, with performance degradation patterns revealing fundamental architectural differences.

Table 4: Inference Speed: Samples per Second by Number of Labels (batch_size = 1, on H100 GPU)

Model Name	1	2	4	8	16	32	64	128	256	512	1024	Average
gliner-bi-edge-v2.0	17.0	27.0	5.05	22.4	17.5	13.9	15.2	12.5	10.8	5.43	3.23	13.64
gliner-bi-edge-v2.0 (pre-computed)	19.3	25.0	28.2	32.6	31.0	32.6	22.2	22.7	22.2	16.9	18.3	24.62
gliner-bi-small-v2.0	12.5	12.8	5.98	11.6	10.6	9.43	6.94	7.35	5.74	3.33	1.60	7.99
gliner-bi-small-v2.0 (pre-computed)	14.7	15.9	14.3	15.3	15.4	15.4	15.6	15.3	15.5	15.7	14.3	15.22
gliner-bi-base-v2.0	8.13	8.62	4.85	8.00	7.52	6.76	5.71	5.21	4.64	3.21	2.30	5.91
gliner-bi-base-v2.0 (pre-computed)	9.52	10.2	9.80	9.95	10.0	9.93	8.93	6.71	9.35	9.71	10.5	9.51
gliner-bi-large-v2.0	3.52	2.53	3.87	3.50	3.66	3.19	1.90	2.46	2.39	1.62	0.87	2.68
gliner-bi-large-v2.0 (pre-computed)	4.37	4.07	4.53	4.54	4.47	3.46	3.85	3.04	2.82	1.84	2.64	3.60
gliner_small-v2.5	10.7	14.6	14.1	13.2	11.9	10.3	7.91	4.26	1.29	0.43	0.14	8.08
gliner_medium-v2.5	7.81	8.51	8.39	7.58	7.12	5.62	4.18	2.19	0.68	0.23	0.07	4.76
gliner_large-v2.5	2.89	3.28	3.29	2.90	2.61	2.33	1.71	1.12	0.31	0.09	0.03	1.87

With pre-computed label embeddings, bi-encoder models maintain near-constant inference speed regardless of label count. GLiNER-bi-edge-v2.0 shows only 5.2% throughput reduction from 1 to 1024 labels (19.3 to 18.3 examples/s), while gliner_small-v2.5 experiences 98.7% degradation (10.7 to 0.14 examples/s). At 1024 labels, the bi-encoder with pre-computation achieves 130 \times higher throughput than the comparable uni-encoder model. Without pre-computation, bi-encoder models still outperform eBERTa-based uni-encoder variants by 23 \times at maximum label count, demonstrating inherent architectural efficiency.

The scalability advantage extends across all model sizes. GLiNER-bi-large-v2.0 maintains 2.64 examples/s at 1024 labels with pre-computation, compared to 0.03 examples/s for gliner_large-v2.5, representing an 88 \times improvement. For production scenarios with 100 entity types, bi-encoder models achieve 5.3 \times higher throughput, translating to 1.96M versus 368K predictions per day on a single H100 GPU.

4 Discussion

4.1 Architectural Trade-offs and Design Rationale

The bi-encoder architecture addresses fundamental scalability limitations of joint encoding approaches through architectural decomposition. But the central question in bi-encoder design is whether decoupling text and label encoding sacrifices the rich cross-modal interactions that joint encoding provides. Our results suggest this concern is largely unfounded for NER and entity linking tasks: GLiNER-bi-large-v2.0 achieves 61.5% on CrossNER compared to 60.9% for gliner_large-v2.5, indicating that late fusion through dot-product scoring suffices for capturing entity-label alignment. We attribute this to the nature of NER, where entity type semantics are relatively self-contained and do not require deep contextual reasoning about the input text during encoding.

The optional cross-attention fusion layer provides a middle ground for scenarios requiring tighter integration. When enabled, bidirectional attention allows entity-aware text representations and text-aware entity representations while preserving the benefits of separate encoding. This design choice reflects a broader principle: architectural modularity enables practitioners to tune the accuracy-efficiency trade-off for their specific deployment constraints.

4.2 The Role of Specialized Encoders

Performance parity with uni-encoder models, despite architectural separation, can be attributed to specialized encoder selection. Sentence transformers optimized for semantic similarity (BGE, MiniLM) provide superior label representation compared to general-purpose encoders. These models are pre-trained on large-scale semantic similarity tasks, making them particularly well-suited for encoding entity type descriptions that must be matched against text spans.

This finding has implications for bi-encoder design more broadly: the choice of label encoder may be as important as the text encoder. Our results show that pairing larger text encoders (ettin-encoder-400m) with correspondingly capable label encoders (bge-base-en-v1.5) yields consistent improvements, while mismatched encoder capacities can lead to bottlenecks. The 194M parameter base variant achieves 98% of the large model’s performance while operating 2.6 \times faster, suggesting that careful encoder selection can achieve favorable accuracy-efficiency trade-offs.

4.3 Domain-Specific Performance Patterns

Analysis of domain-specific results reveals instructive patterns about the strengths and limitations of bi-encoder architectures. Biomedical benchmarks exhibit strong scaling effects, with GENIA_NER improving from 50.1% to 60.1% (20% relative gain) between edge and large variants. This suggests that biomedical entity recognition benefits substantially from increased model capacity, likely due to the complex and specialized vocabulary involved.

In contrast, social media datasets (Broad Tweet Corpus, TweetNER7) show relatively stable performance across model sizes (70.0-72.1% and 36.9-41.7% respectively). Informal text characteristics—non-standard spelling, abbreviations, and context-dependent meanings—may benefit less from increased model capacity and more from domain-specific pre-training or architectural modifications.

The bi-encoder architecture demonstrates particular advantages on datasets where semantic similarity between labels and text is crucial. Superior performance on GENIA_NER (60.1% vs 46.9%), MultiNERD (68.0% vs 50.9%), and Ontonotes (33.3% vs 22.8%) compared to uni-encoder models suggests that dedicated label encoders capture entity type semantics more effectively when label descriptions are informative and well-defined.

Conversely, uni-encoder models maintain advantages on established benchmarks like CoNLL 2003 (65.4% vs 66.5%), possibly because joint encoding better captures the specific annotation conventions and entity boundary patterns learned during pre-training on similar data distributions.

4.4 Scalability and Practical Deployment

The consistent inference speed with pre-computed embeddings transforms large-scale NER deployment from computationally prohibitive to practical. GLiNER-bi-edge-v2.0 shows only 5.2% throughput reduction from 1 to 1024 labels with pre-computed embeddings, while comparable uni-encoder models experience 98.7% degradation. This capability proves particularly valuable for several real-world scenarios.

For biomedical applications that rely on extensive ontologies such as UMLS (4M+ concepts), pre-computing and caching entity embeddings enables practical deployment. Traditional approaches would require encoding millions of labels jointly with each input, rendering real-time inference infeasible. With the bi-encoder architecture, label embeddings can be computed once and stored in efficient vector databases, reducing the problem to text encoding followed by nearest-neighbor search.

Enterprise systems requiring dynamic taxonomy updates also benefit substantially. When new entity types are added or existing ones modified, only the affected label embeddings need recomputation, rather than reprocessing the entire model or invalidating cached computations. This modularity aligns well with production ML systems that must evolve continuously.

For production scenarios with 100 entity types, bi-encoder models achieve $5.3\times$ higher throughput, translating to 1.96M versus 368K predictions per day on a single H100 GPU. This difference can determine whether a system meets latency requirements for user-facing applications or requires costly horizontal scaling.

4.5 Extension to Entity Linking

The bi-encoder architecture's natural extension to entity linking represents a significant practical contribution. Entity linking extends NER by disambiguating mentions to specific knowledge base entities—a task that requires candidate retrieval from potentially millions of entities, followed by disambiguation. The bi-encoder design maps directly to this workflow: pre-computed entity embeddings enable efficient candidate retrieval via approximate nearest neighbor search, while the same scoring mechanism used for NER provides disambiguation.

GLiNER demonstrates this capability through a modular DAG-based pipeline that integrates GLiNER-bi-Encoder as a neural disambiguation component. The framework supports on-the-fly embedding, caching, and database integration for million-scale entity ontologies, addressing a key deployment challenge for knowledge-intensive applications. This unification of NER and entity linking within a single architectural framework simplifies system design and enables end-to-end optimization.

4.6 Limitations and Future Directions

Several limitations warrant discussion. First, performance on highly contextual datasets remains challenging. HarveyNER (10.6-15.0%) and FabNER (22.4-24.3%) show limited improvements with scaling, indicating that certain entity types require deeper contextual reasoning than current architectures provide. These datasets feature entities whose boundaries and types depend heavily on surrounding context—a scenario where joint encoding may provide fundamental advantages.

Second, the fixed maximum span width constraint ($K = 12$ in our experiments) limits recognition of longer entity mentions. It can be sufficient for most NER benchmarks, but specialized domains such as legal or scientific text may contain entity mentions exceeding this limit. While we developed a token-level variant of the bi-encoder architecture, further research is needed to explore its benefits.

Third, while cross-attention fusion is available, our experiments primarily evaluate the base bi-encoder configuration. A systematic study of when and to what extent cross-modal interaction benefits different domains would provide valuable guidance for practitioners.

Future work could explore several promising directions. Contrastive pre-training objectives specifically designed for the bi-encoder NER setting may improve label-text alignment. Integration with retrieval-augmented approaches could enable even larger label vocabularies by dynamically retrieving relevant entity types rather than scoring against all candidates. Finally, extension to multilingual and cross-lingual settings, leveraging multilingual sentence transformers as label encoders, could enable zero-shot NER across languages without parallel training data.

4.7 Broader Implications

The success of bi-encoder architectures for NER contributes to a broader trend in NLP toward modular, decomposable systems. Rather than monolithic models that jointly process all inputs, decomposed architectures enable independent

component optimization, efficient caching and retrieval, and flexible deployment configurations. This paradigm shift has implications beyond NER—similar architectural principles may benefit other structured prediction tasks where input components have separable semantics.

The $130\times$ speedup at scale also highlights the importance of computational efficiency in democratizing NLP capabilities. Models that achieve comparable accuracy with dramatically reduced inference costs enable deployment in resource-constrained environments and reduce the environmental impact of large-scale text processing. As NLP systems are deployed more broadly, such efficiency gains become increasingly important.

5 Conclusion

We presented GLiNER-bi-Encoder, a bi-encoder architecture for named entity recognition that achieves state-of-the-art zero-shot performance (61.5% F1 on CrossNER) while demonstrating superior scalability. The architectural separation of text and label encoding enables $130\times$ faster inference than traditional approaches at 1024 entity types, with near-constant complexity when using pre-computed label embeddings.

Experimental evaluation across 26 NER benchmarks confirms that bi-encoder models match or exceed uni-encoder performance while offering transformative computational advantages. The ability to handle thousands of entity types with sub-second latency addresses a critical limitation in current NER systems, enabling deployment in large-scale production environments previously constrained by computational requirements. Moreover, it enables new use cases, such as entity linking.

The models and implementation are publicly available through the GLiNER library and the Hugging Face Model Hub. This work establishes bi-encoder architectures as a practical solution for scalable information extraction, particularly in domains that require extensive entity taxonomies or real-time processing.

6 Availability

Models are available for use with the Python library GLiNER at <https://github.com/urchade/GLiNER>, and also through the framework UTCA for building pipelines at <https://github.com/Knowledgator/utca>. Models can be downloaded from the Hugging Face repository at <https://huggingface.co/collections/knowledgator/gliner-bi-v2-68ac5880c610ed907cd68a5a>.

Pre-trained models can be downloaded from the Hugging Face repository at: <https://huggingface.co/collections/knowledgator/gliclass-v3-687a2d211b89659da1e3f34a>

7 Acknowledgments

We sincerely thank Urchade Zaratiana, the creator of the original GLiNER-uni-Encoder, and Tom Aarsen, the maintainer of Sentence Transformers.

References

- Olivier Bodenreider. The unified medical language system (umls): integrating biomedical terminology. *Nucleic acids research*, 32(suppl_1):D267–D270, 2004.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.
- Zhi Hong, Logan T. Ward, Kyle Chard, Ben Blaiszik, and Ian T. Foster. Challenges and advances in information extraction from scientific literature: a review. *JOM*, 73:3383 – 3400, 2021. URL <https://api.semanticscholar.org/CorpusID:242354315>.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection, 2018. URL <https://arxiv.org/abs/1708.02002>.
- Ryan McDonald and Fernando Pereira. Identifying gene and protein mentions in text using conditional random fields. *BMC bioinformatics*, 6:1–7, 2005.
- OpenAI, Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, Red Avila, Igor Babuschkin, Suchir Balaji, Valerie Balcom, Paul Baltescu, Haiming Bao, Mohammad Bavarian, Jeff Belgum, Irwan Bello, Jake Berdine, Gabriel

Bernadett-Shapiro, Christopher Berner, Lenny Bogdonoff, Oleg Boiko, Madelaine Boyd, Anna-Luisa Brakman, Greg Brockman, Tim Brooks, Miles Brundage, Kevin Button, Trevor Cai, Rosie Campbell, Andrew Cann, Brittany Carey, Chelsea Carlson, Rory Carmichael, Brooke Chan, Che Chang, Fotis Chantzis, Derek Chen, Sully Chen, Ruby Chen, Jason Chen, Mark Chen, Ben Chess, Chester Cho, Casey Chu, Hyung Won Chung, Dave Cummings, Jeremiah Currier, Yunxing Dai, Cory Decareaux, Thomas Degry, Noah Deutsch, Damien Deville, Arka Dhar, David Dohan, Steve Dowling, Sheila Dunning, Adrien Ecoffet, Atty Eleti, Tyna Eloundou, David Farhi, Liam Fedus, Niko Felix, Simón Posada Fishman, Juston Forte, Isabella Fulford, Leo Gao, Elie Georges, Christian Gibson, Vik Goel, Tarun Gogineni, Gabriel Goh, Rapha Gontijo-Lopes, Jonathan Gordon, Morgan Grafstein, Scott Gray, Ryan Greene, Joshua Gross, Shixiang Shane Gu, Yufei Guo, Chris Hallacy, Jesse Han, Jeff Harris, Yuchen He, Mike Heaton, Johannes Heidecke, Chris Hesse, Alan Hickey, Wade Hickey, Peter Hoeschele, Brandon Houghton, Kenny Hsu, Shengli Hu, Xin Hu, Joost Huizinga, Shantanu Jain, Shawn Jain, Joanne Jang, Angela Jiang, Roger Jiang, Haozhun Jin, Denny Jin, Shino Jomoto, Billie Jonn, Heewoo Jun, Tomer Kaftan, Łukasz Kaiser, Ali Kamali, Ingmar Kanitscheider, Nitish Shirish Keskar, Tabarak Khan, Logan Kilpatrick, Jong Wook Kim, Christina Kim, Yongjik Kim, Jan Hendrik Kirchner, Jamie Kiro, Matt Knight, Daniel Kokotajlo, Łukasz Kondraciuk, Andrew Kondrich, Aris Konstantinidis, Kyle Kosic, Gretchen Krueger, Vishal Kuo, Michael Lampe, Ika Lan, Teddy Lee, Jan Leike, Jade Leung, Daniel Levy, Chak Ming Li, Rachel Lim, Molly Lin, Stephanie Lin, Mateusz Litwin, Theresa Lopez, Ryan Lowe, Patricia Lue, Anna Makanju, Kim Malfacini, Sam Manning, Todor Markov, Yaniv Markovski, Bianca Martin, Katie Mayer, Andrew Mayne, Bob McGrew, Scott Mayer McKinney, Christine McLeavey, Paul McMillan, Jake McNeil, David Medina, Aalok Mehta, Jacob Menick, Luke Metz, Andrey Mishchenko, Pamela Mishkin, Vinnie Monaco, Evan Morikawa, Daniel Mossing, Tong Mu, Mira Murati, Oleg Murk, David Mély, Ashvin Nair, Reiichiro Nakano, Rajeev Nayak, Arvind Neelakantan, Richard Ngo, Hyeonwoo Noh, Long Ouyang, Cullen O'Keefe, Jakub Pachocki, Alex Paino, Joe Palermo, Ashley Pantuliano, Giambattista Parascandolo, Joel Parish, Emy Parparita, Alex Passos, Mikhail Pavlov, Andrew Peng, Adam Perelman, Filipe de Avila Belbute Peres, Michael Petrov, Henrique Ponde de Oliveira Pinto, Michael Pokorny, Michelle Pokrass, Vitchyr H. Pong, Tolly Powell, Alethea Power, Boris Power, Elizabeth Proehl, Raul Puri, Alec Radford, Jack Rae, Aditya Ramesh, Cameron Raymond, Francis Real, Kendra Rimbach, Carl Ross, Bob Rotsted, Henri Roussez, Nick Ryder, Mario Saltarelli, Ted Sanders, Shibani Santurkar, Girish Sastry, Heather Schmidt, David Schnurr, John Schulman, Daniel Selsam, Kyla Sheppard, Toki Sherbakov, Jessica Shieh, Sarah Shoker, Pranav Shyam, Szymon Sidor, Eric Sigler, Maddie Simens, Jordan Sitkin, Katarina Slama, Ian Sohl, Benjamin Sokolowsky, Yang Song, Natalie Staudacher, Felipe Petroski Such, Natalie Summers, Ilya Sutskever, Jie Tang, Nikolas Tezak, Madeleine B. Thompson, Phil Tillet, Amin Toootoonchian, Elizabeth Tseng, Preston Tuggle, Nick Turley, Jerry Tworek, Juan Felipe Cerón Uribe, Andrea Vallone, Arun Vijayvergiya, Chelsea Voss, Carroll Wainwright, Justin Jay Wang, Alvin Wang, Ben Wang, Jonathan Ward, Jason Wei, CJ Weinmann, Akila Welihinda, Peter Welinder, Jiayi Weng, Lilian Weng, Matt Wiethoff, Dave Willner, Clemens Winter, Samuel Wolrich, Hannah Wong, Lauren Workman, Sherwin Wu, Jeff Wu, Michael Wu, Kai Xiao, Tao Xu, Sarah Yoo, Kevin Yu, Qiming Yuan, Wojciech Zaremba, Rowan Zellers, Chong Zhang, Marvin Zhang, Shengjia Zhao, Tianhao Zheng, Juntang Zhuang, William Zhuk, and Barret Zoph. Gpt-4 technical report, 2024.

Georgios Petasis, Frantz Vichot, Francis Wolinski, Georgios Paliouras, Vangelis Karkaletsis, and Constantine D Spyropoulos. Using machine learning to maintain rule-based named-entity recognition and classification systems. In *proceedings of the 39th annual meeting of the association for computational linguistics*, pages 426–433, 2001.

Kristie Seymore, Andrew McCallum, Roni Rosenfeld, et al. Learning hidden markov model structure for information extraction. In *AAAI-99 workshop on machine learning for information extraction*, pages 37–42, 1999.

Lucia Siciliani, Eleonora Ghizzota, Pierpaolo Basile, and Pasquale Lops. Oie4pa: open information extraction for the public administration. *J. Intell. Inf. Syst.*, 62:273–294, 2023. URL <https://api.semanticscholar.org/CorpusID:262178371>.

Matyáš Skalický, Stepán Simska, Michal Uřičář, and Milan Šulc. Business document information extraction: Towards practical benchmarks. *ArXiv*, abs/2206.11229, 2022. URL <https://api.semanticscholar.org/CorpusID:249926391>.

Orion Weller, Kathryn Ricci, Marc Marone, Antoine Chaffin, Dawn Lawrie, and Benjamin Van Durme. Seq vs seq: An open suite of paired encoders and decoders, 2025. URL <https://arxiv.org/abs/2507.11412>.

Shitao Xiao, Zheng Liu, Peitian Zhang, and Niklas Muennighoff. C-pack: Packaged resources to advance general chinese embedding, 2023.

Anthony Yazdani, Ihor Stepanov, and Douglas Teodoro. Gliner-biomed: A suite of efficient models for open biomedical named entity recognition. *arXiv preprint arXiv:2504.00676*, 2024.

Urchade Zaratiana, Nadi Tomeh, Pierre Holat, and Thierry Charnois. Gliner: Generalist model for named entity recognition using bidirectional transformer, 2023.